

Vision-text representational similarity across layers
Qwen2-VL-7B-Instruct, 500 pairs each, centered per-pair cosines

