

Vision-text similarity across layers — MISMATCHED pairs
Qwen2-VL-7B-Instruct, 500 pairs, layer-0 ablation

